GEMM leaderboard

One kernel, every workload. Each row is a **single run**: all four numbers come from the same kernel, so a kernel has to be fast at the square GEMM, at N=1 (matrix-vector, the language model's case), at N=32 *and* on TinyStories at once. Optimizing for just one of them won't do. You should pick the run shown here on your page (☆ next to the run); until you do, it's your run with the highest tok/s.

Click a column to rank by it. 0 job(s) in the queue.

rank	roll	square 768^3 (ms) ▼	N=1 total (ms)	N=32 total (ms)	TinyStories tok/s
1	2023111030	0.986	3.085	3.104	117.35
2	2023111014	1.654	3.549	3.712	98.27
3	2024101012	2.466	10.395	8.222	45.58
4	2023113009	3.820	8.466	9.364	53.18
5	2023112023	3.861	6.725	12.569	63.17
6	2023112003	4.002	3.208	3.581	117.57
7	2023111019	6.451	3.212	4.384	118.72
8	2025201015	7.050	6.839	7.742	60.79
9	2023102030	7.081	6.680	38.015	66.84
10	2023101096	7.166	3.237	5.393	106.61
11	2023111003	7.447	6.765	5.322	64.85
12	2025121012	7.463	3.242	5.629	115.38
13	2024101002	7.703	37.590	5.332	15.61
14	2023113005	7.796	3.753	5.196	105.46
15	2024121004	8.140	6.670	5.224	69.96
16	2023101080	8.713	27.208	24.759	20.52
17	2023111015	10.266	7.596	7.884	54.09
18	2024112019	10.986	4.091	8.427	97.22
19	2023112005	12.395	3.758	8.648	104.25
20	2024102051	12.745	14.200	14.198	35.47
21	2023111004	13.024	6.806	8.842	61.64

22	2025121003	25.424	16.652	33.771	31.44
23	2023115018	33.080	4.359	17.867	86.12
24	2025901017	46.798	61.352	72.118	9.30
25	2026203010	159.469	111.895	169.120	5.44
26	2024112022	1514.686	88.123	1099.446	6.84