GEMM leaderboard

One kernel, every workload. Each row is a **single run**: all four numbers come from the same kernel, so a kernel has to be fast at the square GEMM, at N=1 (matrix-vector, the language model's case), at N=32 *and* on TinyStories at once. Optimizing for just one of them won't do. You should pick the run shown here on your page (☆ next to the run); until you do, it's your run with the highest tok/s.

Click a column to rank by it. 0 job(s) in the queue.

rank	roll	square 768^3 (ms)	N=1 total (ms)	N=32 total (ms)	TinyStories tok/s ▼
1	2023111019	6.451	3.212	4.384	118.72
2	2023111030	0.986	3.085	3.104	117.35
3	2025121012	7.463	3.242	5.629	115.38
4	2023112003	2.146	3.200	3.217	114.66
5	2023101096	7.166	3.237	5.393	106.61
6	2023113005	7.796	3.753	5.196	105.46
7	2023112005	12.395	3.758	8.648	104.25
8	2023111014	1.654	3.549	3.712	98.27
9	2024112019	10.986	4.091	8.427	97.22
10	2023115018	33.080	4.359	17.867	86.12
11	2025201015	13.484	6.480	60.913	74.26
12	2024121004	8.140	6.670	5.224	69.96
13	2023111003	7.447	6.765	5.322	64.85
14	2023112023	10267.753	6.948	5769.922	62.65
15	2024101012	2.466	10.395	8.222	45.58
16	2025121003	25.424	16.652	33.771	31.44
17	2024102051	24.419	22.445	31.168	24.57
18	2023101080	8.713	27.208	24.759	20.52
19	2024101002	7.703	37.590	5.332	15.61
20	2023111015	28.674	44.124	46.120	12.90
21	2023102030	7.207	46.983	38.603	11.50

22	2025901017	46.798	61.352	72.118	9.30
23	2024112022	1514.686	88.123	1099.446	6.84
24	2023111004	1083.006	100.696	998.629	5.97
25	2026203010	159.469	111.895	169.120	5.44